

Bayesian Modelling - Homework 1

Sebastian Kimm-Friedenberg

2022-11-02

Load packages:

```
library(bayesm)
library(data.table)
library(knitr)
library(flextable)
```

```
## Warning: Paket 'flextable' wurde unter R Version 4.2.2 erstellt
```

```
library(tidyverse)
set.seed(60323)

library(compiler)
runireg_c=cmpfun(runireg) # a bit faster
```

Task 1

```
nobs <- c(5,10,100,1000,10000,20000, 30000) # different ns

# matrices in which results are stored
b_b1 <- matrix(0, length(nobs), 3) # m rows, 3 columns
sd_b1 <- matrix(0, length(nobs), 3)
b_ols <- b_b1
sd_ols <- sd_b1

iteration <- 0

for (n in nobs){
  print(n)
  iteration <- iteration + 1
  # set up DGP
  #n <- 5
  X <- cbind(rep(1,n),runif(n),runif(n))
  beta <- c(-1,3,6)
  y <- (X %*% beta + rnorm(n, mean = 0, sd = 1))# set up y

  # set up sampling
  R <- 80000 # as in the lecture
  Mcmc1 <- list(R=R,keep=1, nprint = 0)
```

```

Prior <- list(nu=0, ssq=0) # beta = 0 as default

# run regressions
## bayes
out_b1 <- runireg_c(Data=list(y=y,X=X), Prior = Prior, Mcmc=Mcmc1)
b_draw1 <- out_b1$betadraw
b_b1[iteration,] <- apply(b_draw1, 2, mean) # get mean for each coeff.
sd_b1[iteration,] <- apply(b_draw1, 2, sd) # standard deviation

## ols
out_ols <- lm(y ~ X[, -1]) # without intercept in x, is added automatically
summary(out_ols)
b_ols[iteration,] <- matrix(out_ols$coefficients, 1, 3)
sd_ols[iteration,] <- matrix(sqrt(diag(vcov(out_ols))), 1, 3)
}

```

The previous code computed the outcomes for the Bayes and OLS estimator along with their standard deviations. The results are presented below:

```

## Warning: Warning: fonts used in 'flextable' are ignored because the 'pdflatex'
## engine is used and not 'xelatex' or 'lualatex'. You can avoid this warning
## by using the 'set_flextable_defaults(fonts_ignore=TRUE)' command or use a
## compatible engine by defining 'latex_engine: xelatex' in the YAML header of the
## R Markdown document.

```

| n | b1 | sd_b1 | b2 | sd_b2 | b3 | sd_b3 | ols1 | sd_ols1 | ols2 | sd_ols2 | ols3 | sd_ols3 |
|--------|---------|--------|--------|--------|--------|--------|---------|---------|--------|---------|--------|---------|
| 5 | 0.1272 | 1.2480 | 1.2040 | 1.4839 | 6.5530 | 1.2107 | 0.0478 | 1.4324 | 1.2464 | 1.7047 | 6.6661 | 1.3786 |
| 10 | -0.0683 | 1.0388 | 3.2107 | 1.0363 | 4.6992 | 1.0713 | -0.2166 | 1.0908 | 3.3151 | 1.0768 | 4.8368 | 1.1239 |
| 100 | -1.4613 | 0.2309 | 3.3252 | 0.3332 | 6.6284 | 0.3038 | -1.4683 | 0.2325 | 3.3301 | 0.3341 | 6.6370 | 0.3059 |
| 1,000 | -1.0997 | 0.0838 | 3.0957 | 0.1071 | 6.0799 | 0.1113 | -1.1002 | 0.0838 | 3.0955 | 0.1074 | 6.0812 | 0.1104 |
| 10,000 | -1.0147 | 0.0263 | 3.0308 | 0.0344 | 5.9962 | 0.0345 | -1.0149 | 0.0263 | 3.0310 | 0.0345 | 5.9963 | 0.0344 |
| 20,000 | -1.0159 | 0.0186 | 3.0272 | 0.0245 | 5.9974 | 0.0244 | -1.0159 | 0.0186 | 3.0271 | 0.0245 | 5.9974 | 0.0244 |
| 30,000 | -1.0188 | 0.0153 | 3.0090 | 0.0199 | 6.0176 | 0.0200 | -1.0188 | 0.0153 | 3.0090 | 0.0199 | 6.0175 | 0.0199 |

We can see that with increasing sample size both estimators converge to the true values of $\beta = (-1, 3, 6)$. The Bayesian estimator constantly has a smaller standard deviation than the OLS estimator, implying the Bayesian being more efficient.

Task 2

```

nobs <- 10000 # medium sized observation size
noby <- n # faster then changing all "n" to "nobs"
# matrices in which results are stored
b_b <- matrix(0, length(nobs), 3) # m rows, 3 columns
sd_b <- matrix(0, length(nobs), 3)

```

```

b_ols <- b_b
sd_ols <- b_b

# set up DGP
#n <- 5
hilmf= runif(nobs)
X <- cbind(rep(1,n),hilmf, hilmf)
beta <- c(-1,3,6)
y <- (X %*% beta + rnorm(n, mean = 0, sd = 1))# set up y

# set up sampling
R <- 80000 # as in the lecture
Mcmc1 <- list(R=R,keep=1, nprint = 0)
Prior <- list(nu=0, ssq=0) # beta = 0 as default

# run regressions
## bayes
out_b <- runireg_c(Data=list(y=y,X=X), Prior = Prior,Mcmc=Mcmc1)
b_draw <-out_b$betadraw
b_b[1,] <- apply(b_draw, 2, mean) # get mean for each coeff.
sd_b[1,] <- apply(b_draw, 2, sd) # standard deviation

## ols
out_ols <- lm(y ~ X[,-1]) # without intercept in x, is added automatically
summary(out_ols)
b_ols[1,] <- matrix(out_ols$coefficients, 1, 3)
sd_ols[1,] <- matrix(sqrt(diag(vcov(out_ols))), 1, 3)

res <- data.frame(n = nobs, b1 = b_b[,1], sd_b1 = sd_b[,1],
                  b2 = b_b[,2], sd_b2 = sd_b[,2],
                  b3 = b_b[,3], sd_b3 = sd_b[,3],
                  ols1 = b_ols[,1], sd_ols1 = sd_ols[,1],
                  ols2 = b_ols[,2], sd_ols2 = sd_ols[,2],
                  ols3 = b_ols[,3], sd_ols3 = sd_ols[,3])

round(res,4) %>% regularartable() %>% autofit() %>% fit_to_width(max_width = 6.5)

```

```

## Warning: Warning: fonts used in 'flextable' are ignored because the 'pdflatex'
## engine is used and not 'xelatex' or 'lualatex'. You can avoid this warning
## by using the 'set_flextable_defaults(fonts_ignore=TRUE)' command or use a
## compatible engine by defining 'latex_engine: xelatex' in the YAML header of the
## R Markdown document.

```

| n | b1 | sd_b1 | b2 | sd_b2 | b3 | sd_b3 | ols1 | sd_ols1 | ols2 | sd_ols2 | ols3 | sd_ols3 |
|--------|---------|--------|--------|-------|--------|-------|---------|---------|--------|---------|------|---------|
| 10,000 | -1.0066 | 0.0116 | 4.5245 | 7.071 | 4.4935 | 7.071 | -1.0067 | 0.0116 | 9.0181 | 0.0202 | | |

```

kable(summary(b_draw), digits = 4)

```

```

Summary of Posterior Marginal Distributions Moments mean std dev num se rel eff sam size
1 -1.0 0.012 0.00004 0.87 72000 2 4.5 7.079 0.02579 0.96 72000 3 4.5 7.079 0.02580 0.96 72000

```

Quantiles 2.5% 5% 50% 95% 97.5% 1 -1.0 -1.0 -1.0 -0.99 -0.98 2 -9.3 -7.2 4.5 16.17 18.42 3 -9.4 -7.1 4.5 16.19
18.32 based on 72000 valid draws (burn-in=8000)

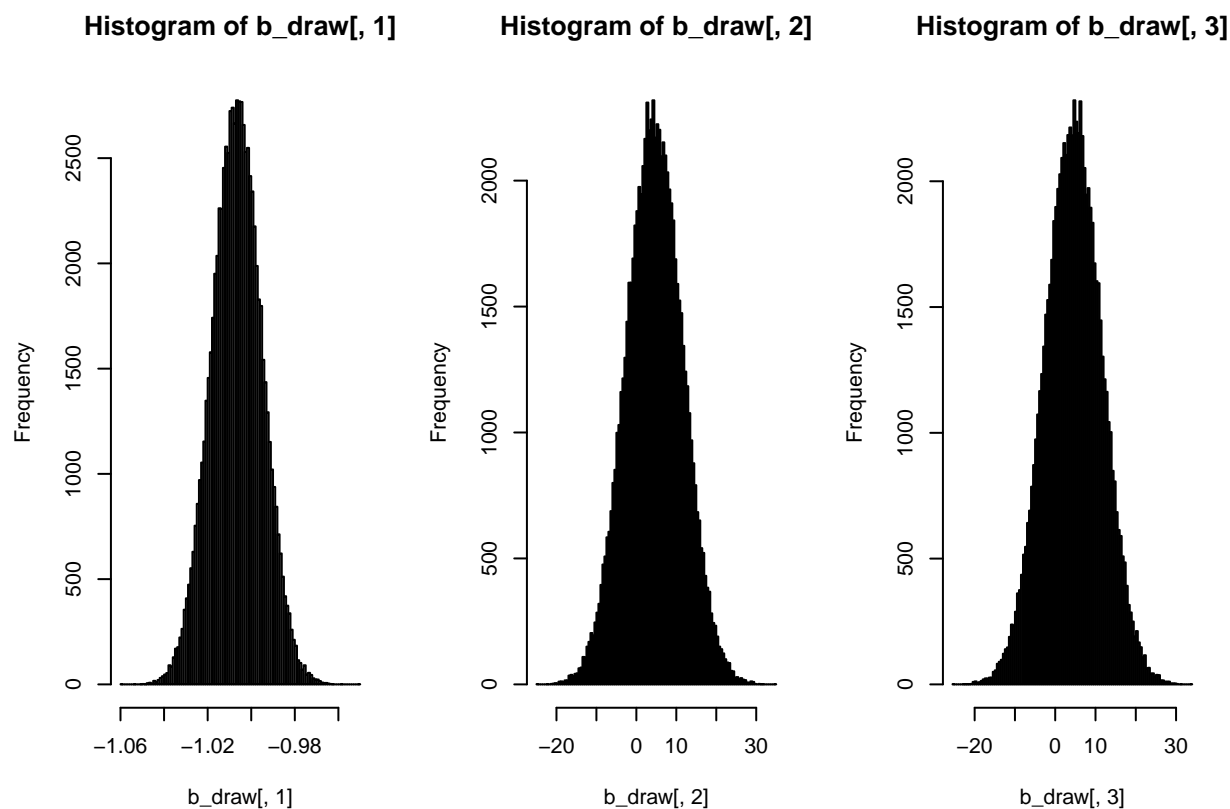
| | mean | std | dev | num | se | rel | eff | size |
|--|---------|--------|---------|------|----|-----|-----|-------|
| | -1.0066 | 0.0116 | 0.00000 | 8744 | | | | 72000 |
| | 4.5339 | 7.0793 | 0.02580 | 9556 | | | | 72000 |
| | 4.4841 | 7.0793 | 0.02580 | 9560 | | | | 72000 |

```
kable(t(apply(b_draw,2,quantile)), digits = 4)
```

| | 0% | 25% | 50% | 75% | 100% |
|--|----------|---------|---------|---------|---------|
| | -1.0594 | -1.0144 | -1.0066 | -0.9989 | -0.9506 |
| | -24.7829 | -0.2485 | 4.4883 | 9.2924 | 34.5134 |
| | -25.4765 | -0.2735 | 4.5263 | 9.2687 | 33.7629 |

The `lm` command, i.e. the OLS estimation, automatically omits one of the identical variables. Due to multicollinearity, the matrix $X'X$ is not invertible and the OLS estimator would not be feasible. In contrast, the Bayesian posterior distribution is computed for both variables and it is almost identical. The third variable has a slightly smaller mean and seems to be a bit more skewed to the right, according to the quantiles.

Plotting the three variables confirms that the posterior distribution of 2 and 3 is almost identical.



Task 3

```
nobs <- c(5,10,100,1000,10000,20000, 30000) # different ns

# matrices in which results are stored
b_b3 <- matrix(0, length(nobs), 3) # m rows, 3 columns
sd_b3 <- matrix(0, length(nobs), 3)
b_pro <- b_b3
sd_pro <- sd_b3

iteration <- 0

for (n in nobs){
  print(n)
  iteration <- iteration + 1
  # set up DGP
  #n <- 5
  X <- cbind(rep(1,n),runif(n),runif(n))
  beta <- c(-1,3,6)
  y <- (X %*% beta + rnorm(n, mean = 0, sd = 1))# set up y
  y <- ifelse(y<0,0,1)

  # set up sampling
  R <- 80000 # as in the lecture
```

```

Mcmc1 <- list(R=R,keep=1, nprint = 0)
#Prior <- list(nu=0, ssq=0) # beta = 0 as default

# run regressions
## bayes
out_b3 <- runireg_c(Data=list(y=y,X=X),Mcmc=Mcmc1)
b_draw3 <-out_b3$betadraw
b_b3[iteration,] <- apply(b_draw3, 2, mean) # get mean for each coeff.
sd_b3[iteration,] <- apply(b_draw3, 2, sd) # standard deviation

## ols
out_pro <- glm(y ~ X[,-1], family=binomial(link="probit")) # without intercept in x, is added automat
summary(out_pro)
b_pro[iteration,] <- matrix(out_pro$coefficients, 1, 3)
sd_pro[iteration,] <- matrix(sqrt(diag(vcov(out_pro))), 1, 3)
}

```

```

## Warning: glm.fit: Angepasste Wahrscheinlichkeiten mit numerischem Wert 0 oder 1
## aufgetreten

```

```

## Warning: glm.fit: Angepasste Wahrscheinlichkeiten mit numerischem Wert 0 oder 1
## aufgetreten

```

```

## Warning: glm.fit: Angepasste Wahrscheinlichkeiten mit numerischem Wert 0 oder 1
## aufgetreten

```

Below, the first table shows the results of the Bayes estimator and the probit estimator for the probit data. The second shows the comparison of the Bayes estimator performance for linear and probit data. Clearly, the probit model converges (slowly) to the true parameters whereas the Bayesian model is consistently wrong, i.e. unable to recover the true parameters -1, 3, 6. Regarding the trade-off between sample size and having access to the linear observations or only to the probit observations, the access to linear observations should always be preferred. In the linear case, the estimator converges relatively quickly whereas in the probit case it quite far away from the true result regardless of the sample size.

```

## Warning: Warning: fonts used in 'flectable' are ignored because the 'pdflatex'
## engine is used and not 'xelatex' or 'lualatex'. You can avoid this warning
## by using the 'set_flectable_defaults(fonts_ignore=TRUE)' command or use a
## compatible engine by defining 'latex_engine: xelatex' in the YAML header of the
## R Markdown document.

```

| n | b1 | sd_b1 | b2 | sd_b2 | b3 | sd_b3 | pro1 | sd_pro1 | pro2 | sd_pro2 | pro3 | sd_pro3 |
|--------|--------|--------|--------|--------|---------|--------|---------|-------------|--------|-------------|--------|--------------|
| 5 | 0.9822 | 0.0543 | 0.0249 | 0.1080 | -0.0096 | 0.2255 | 6.5528 | 24,074.9174 | 0.0000 | 54,796.1802 | 0.0000 | 120,125.7980 |
| 10 | 0.9925 | 0.0262 | 0.0073 | 0.0319 | 0.0059 | 0.0296 | 6.5528 | 15,218.3276 | 0.0000 | 18,556.5465 | 0.0000 | 17,132.8289 |
| 100 | 0.6354 | 0.0617 | 0.2503 | 0.0786 | 0.3367 | 0.0823 | -3.2218 | 1.3506 | 7.2138 | 2.7046 | 9.7542 | 4.2229 |
| 1,000 | 0.7555 | 0.0169 | 0.1788 | 0.0224 | 0.2131 | 0.0218 | -0.7046 | 0.2176 | 2.9741 | 0.4333 | 5.1495 | 0.7671 |
| 10,000 | 0.7311 | 0.0056 | 0.1658 | 0.0074 | 0.2592 | 0.0074 | -1.0011 | 0.0743 | 2.7890 | 0.1313 | 6.0738 | 0.2566 |
| 20,000 | 0.7385 | 0.0039 | 0.1754 | 0.0050 | 0.2414 | 0.0051 | -1.0801 | 0.0550 | 3.2650 | 0.1065 | 6.0528 | 0.1877 |

| n | b1 | sd_b1 | b2 | sd_b2 | b3 | sd_b3 | pro1 | sd_pro1 | pro2 | sd_pro2 | pro3 | sd_pro3 |
|--------|--------|--------|--------|--------|--------|--------|---------|---------|--------|---------|--------|---------|
| 30,000 | 0.7370 | 0.0032 | 0.1741 | 0.0041 | 0.2446 | 0.0041 | -1.0379 | 0.0436 | 3.1232 | 0.0828 | 6.0727 | 0.1529 |

```
## Warning: Warning: fonts used in 'flextable' are ignored because the 'pdflatex'
## engine is used and not 'xelatex' or 'lualatex'. You can avoid this warning
## by using the 'set_flextable_defaults(fonts_ignore=TRUE)' command or use a
## compatible engine by defining 'latex_engine: xelatex' in the YAML header of the
## R Markdown document.
```

| n | b1.1 | sd_b1.1 | b2.1 | sd_b2.1 | b3.1 | sd_b3.1 | b1.3 | sd_b1.3 | b2.3 | sd_b2.3 | b3.3 | sd_b3.3 |
|--------|---------|---------|--------|---------|--------|---------|--------|---------|--------|---------|---------|---------|
| 5 | 0.1272 | 1.2480 | 1.2040 | 1.4839 | 6.5530 | 1.2107 | 0.9822 | 0.0543 | 0.0249 | 0.1080 | -0.0096 | 0.2255 |
| 10 | -0.0683 | 1.0388 | 3.2107 | 1.0363 | 4.6992 | 1.0713 | 0.9925 | 0.0262 | 0.0073 | 0.0319 | 0.0059 | 0.0296 |
| 100 | -1.4613 | 0.2309 | 3.3252 | 0.3332 | 6.6284 | 0.3038 | 0.6354 | 0.0617 | 0.2503 | 0.0786 | 0.3367 | 0.0823 |
| 1,000 | -1.0997 | 0.0838 | 3.0957 | 0.1071 | 6.0799 | 0.1113 | 0.7555 | 0.0169 | 0.1788 | 0.0224 | 0.2131 | 0.0218 |
| 10,000 | -1.0147 | 0.0263 | 3.0308 | 0.0344 | 5.9962 | 0.0345 | 0.7311 | 0.0056 | 0.1658 | 0.0074 | 0.2592 | 0.0074 |
| 20,000 | -1.0159 | 0.0186 | 3.0272 | 0.0245 | 5.9974 | 0.0244 | 0.7385 | 0.0039 | 0.1754 | 0.0050 | 0.2414 | 0.0051 |
| 30,000 | -1.0188 | 0.0153 | 3.0090 | 0.0199 | 6.0176 | 0.0200 | 0.7370 | 0.0032 | 0.1741 | 0.0041 | 0.2446 | 0.0041 |